Net Gains: Systematic Phishing Page Differentiation by JavaScript Behavior

Anonymous authors

Abstract

Phishing remains one of the greatest security challenges facing users, while the Anti-Phishing Work Group reported over one million phishing pages in the last year alone. Phishers achieve this scale by using phishing kits — ready-to-deploy phishing websites — to rapidly deploy phishing campaigns with similar data exfiltration, evasion, or mimicry techniques. In contrast, researchers and defenders continue to fight phishing on a page-by-page basis, and they still rely on manual analysis to recognize static features for kit identification.

This paper aims to aid researchers and analysts in studying phishing by automatically clustering pages in large datasets via the page's browser API usage. In most cases, our techniques separate pages by the underlying phishing kits from which they originate. Our system has an accuracy of 93% on a ground-truth dataset of pages and kits collected from the wild. With a curated mapping of techniques to browser APIs and over 500,000 pages in which we identify 11,377 clusters, we explore what techniques are universal, widespread across kits, or kit-specific. We find UI interactivity and basic fingerprinting (User-agent, Cookies, and referer) to be the most universal techniques, present in 90% and 80% of the clusters, respectively. On the other hand, mouse detection via the browser's mouse API is among the rarest behaviors, despite being used in a deployment of a 7-year-old open-source phishing kit. Overall, we show how defenders can leverage the increasing JavaScript complexity of phishing pages against adversaries, identifying phishing pages originating from the same kits, thus enabling defenders to work at more scalable abstraction levels.

1 Introduction

Web-based phishing attacks, where a webpage, through mimicking or urgency, tricks the user into submitting personal information to an attacker, have been increasing for the last 5 years [13]. While varying in delivery vectors, sent through email, SMS, or QR codes, phishing attacks remain successful in compromising individual accounts and enterprises.

In 2021, the US Cybersecurity and Infrastructure Security Agency (CISA) warned of phishing campaigns targeting Nongovernmental organizations (NGO) and government workers that smuggled the HTML and JavaScript of the phishing page through an attachment [6,27]. Once an actor steals credentials, they sell them on illicit markets or leverage them to leak more data from the target. For example 2024, with the Multi-State Information Sharing and Analysis Center, CISA reported that compromised former employee credentials were being used to access internal networks within the US government [7].

Phishing kits, ready-to-deploy software packages sold at illicit markets for launching phishing attacks, have lowered the barrier of entry for malicious actors. Sellers often market phishing kits as bundles of quality-of-life features for attackers, such as built-in evasions from automated crawlers, exfiltration to Telegram channels, and obfuscation [44, 58]. Deploying these kits can be as easy as uploading them to free hosting providers and mass sending multiple links that exfiltrate the credentials to the same endpoints.

One of these selling points, evasions, extends a phishing page's lifetime (delta time between deployment and discovery) and increases the number of victims visiting the page without a browser warning. While server-side logic of kits can only make assumptions about the system based on the IP address and User agent, through API calls to the browser, client-side JavaScript code can query the user's system for CPU core counts, memory overhead, request user interaction, or call out to a 3rd party bot detection like CloudFlare [62, 64]. In the end, the JavaScript logic in a kit can range from a simple user-agent-based redirection to an AES-encrypted script that dynamically decrypts itself, identifies the browser through a series of API calls, and renders the page after confirming the victim is using a mobile device.

We leverage the complexity of the page's JavaScript logic to cluster the page and further annotate the clusters with what phishing techniques they employ. This paper aims to aid researchers and analysts by automatically differentiating groups of phishing pages based on the underlying kit. Crawling phishing pages from APWG, PhishTank, OpenPhish, URLScan, and

PhishingDB for *545*, we collect 533,514 pages with sophisticated enough JavaScript logic (8 Browser APIs executed) and 7,273 Zip files as potential phishing kits for these websites. Through SHA256 deduplication, and manual examination (Described in Section-XX), we distil this 521 phishing kit families deployed across 4,448. Evaluating the crawler over this dataset of URLs and kits, we find that clustering over the set of browser APIs executed using hierarchical clustering yields a completeness score of XXX and a homogeneity score of XXX. We then turn the clustering to 533,514 pages, which we group into *11,377* clusters. Figure 1 shows how the volume of pages/week can be by an order of magnitude when viewed as active clusters/week. We then use a predefined mapping of browser APIs to phishing techniques to better understand how widespread these practices are, while controlling for the mass-deployment of a single kit. For example, Client-Side IP checks occur in **19,869** pages; however, this totals to **504** clusters, and **23**% of the pages come from a single cluster, which consists of Facebook business account phishing pages. Overall, we offer the following contributions:

C1. We demonstrate that *browser API usage alone* is sufficient to isolate and distinguish known phishing kits. With a ground truth dataset of 526 phishing kits deployed across 4,448 pages, we achieve accuracy metrics of a 93% Fowlkes-Mallows score and an 86% V-measure in the clusters of these pages relative to their kits used.

C2. We experimentally demonstrate optimal parameters for this kind of clustering. We argue that browser APIs common on the web (DOM APIs, property reads, etc.) serve as a valuable identifier for identifying kits, as they signal the kit developer's choices.

C3. We digest 434,050 phishing feeds into *11,377* clusters, which allows us to observe that most clusters have short lives, and target a single brand, suggesting that mass phishing campaigns gravitate towards low-quality, non-modular, and cheap or free kits.

C4. We highlight how widespread different techniques are in the ecosystem by annotating dynamic traces of pages within a cluster and propagating the behavior to the cluster as a whole. When annotating all the clusters for the different techniques used, we find that UI-interactivity and basic fingerprinting are near-universal nowadays. At the same time, mouse detection via browser APIs, Cloudflare Turnstile embedding, and dynamic script creation are still relatively rare. We also note that compared to prior work, we observe a decreased ratio of pages that employ pop-ups as a form of bot detection.

C5. We release the dataset of ground truth labeled URLs we collected and all the clusters we observed.

This work stands apart from prior research, as we automatically differentiate pages via all their browser APIs executed instead of static features or pre-trained classifiers. In doing so, we achieve a much higher accuracy than prior work, of F1=39.54% from [16]. Leveraging prior work, we build a
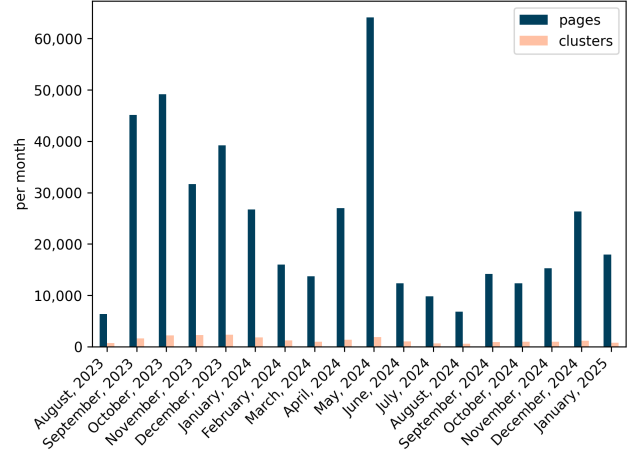

Figure 1: Comparison between monthly pages observed vs the monthly clusters observed based on dynamic behaviors. We observe a drastic, non-linear reduction of phenomena that need to be investigated monthly.

mapping of different evasion techniques (Fingerprinting, Obfuscation, CloudFlare Turnstile, Client-side IP check, etc.) that a phishing page may employ and provide an up-to-date and mass-deployment-resilient description of how widespread these are.

2 Background

In this section, we provide a background on current developments in phishing as a phenomenon and adversarial JavaScript techniques, as well as an overview of the parties involved in the phishing ecosystem.

2.1 The Phenomenon of Phishing

Phishing is a form of social engineering where an adversary pretends to be a trusted entity to steal a user's credentials or gain access to a specific machine, network, or account to which the user has access. While delivery mechanisms vary, most phishing eventually leads to a webpage that requests some personal information (usually credentials) from the user. Because some legitimate websites use web fingerprints as a secondary authentication vector, phishers now also use browser fingerprinting APIs to identify real users and exfiltrate fingerprints to pair with stolen credentials [52].

Phishing is ever-evolving and still growing in prevalence. Groups that track phishing saw an increase in phishing domains in the 2020s. The most popular target sectors vary each year, but include software-as-a-service and webmail services (Q3 2020), financial services (Q3 of 2021), and social media (2024) [1]. In Q3 2024 alone, the Anti-Phishing Working Group (APWG) reported 900,000 phishing attacks.

As more enterprises and researchers study and combat phishing, phishers respond with new countermeasures to prevent automated crawling and phishing detection, collectively called "cloaking." If a phishing page determines the client to be not a viable victim (e.g., a crawling bot, not in a specific country, etc.), it take actions to not serve real phishing content. The page may simply halt with an empty DOM or redirect to a benign page, a long-dead phishing page, or an affiliate marketing page.

Cloaking techniques can be broadly classified into server-side and client-side techniques [44]. Server-side techniques are stealthier, but they rely on limited information about the client. Most server-side techniques rely on precompiled deny-lists or allow-lists of IP addresses, user agents, or referers (the page from which the link is visited as identified by an HTTP header). Although server-side code is usually inaccessible, researchers can still analyze it when phishers leave behind predeployment assets, usually in zip bundles called "phishing kits."

Client-side techniques allow for richer evasion strategies, but they are also more detectable. Phishing pages use browser APIs to trigger permission pop-ups to identify crawling browsers, which often cannot interact with the whole browser UI. Because cloaking is technically very similar to legitimate bot-detection and abuse prevention, phishers also use CAPTCHAs and click-through pages as a form of client-side cloaking. Recently, phishing pages have used Cloudflare Turnstile, a popular widget for abuse prevention, to identify automated browsers. Figure 2 shows a phishing page using a Cloudflare Turnstile when we crawled it.

Even the URL features of a phishing link contain techniques that have evolved to respond to anti-phishing research. Phishing pages frequently use URL shorteners (public or private) to obfuscate the final destination, landing pages requiring a user to follow hyperlinks to the actual page, and free web hosting with trustworthy top-level domains (TLDs).
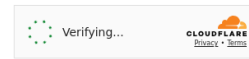
2.2 Phishing as an ecosystem

Phishing is a logistical and technical challenge because a phisher must develop an effective phishing page with cloaking; find robust hosting for it; entice a victim to browse to it through SMS, email, or social media; exfiltrate the phished data; then monetize the stolen credentials. This technical and logistical complexity, combined with interest from potential phishers, has created an underground economy to facilitate each step.

Phishing facilitators sell bundles of customizable or ready-to-deploy phishing pages known as "phishing kits." They vary in features, sophistication, and cost. Phishing-as-a-service providers offer phishing kits combined with hosting services — essentially turnkey phishing systems. Both products lower barriers to entry. Prior work has shown that phishing kits may steal credentials from their customers' kit deployments

Figure 2: Example of a phishing page in our dataset that embedded Cloudflare Turnstile verification on a non-Cloudflare domain

[20, 38], adapt or "borrow" features from other kits [28], and they are sometimes tied to specific actors [10]. Phishing systems may store credentials on the same server, risking loss when the page is inevitably taken down, but a more common practice is to send them to the phisher over instant messaging channels.

Credential sales markets simplify monetization. The credential sales part of the ecosystem has also adapted to modern MFA/2FA practices. With prior work showing that a browser fingerprint is enough to trick online services to trigger an MFA bypass [35], and has an increasing effect on the costs of stolen credentials [52].

2.3 JavaScript

Originally meant as a way of adding interactivity to webpages. The JavaScript ecosystem has evolved to allow varying low-level features to webpages through Browser APIs. HTML DOM APIs enable developers to modify page appearance, while LocalStorage and IndexDB allow write access to the browser's internal storage buckets; meanwhile, the File System API can allow access to the user's real machine's storage. Browser APIs can be function calls (or constructors), property reads, and property writes. Most of the privileged functionality comes from function calls.

The dynamic nature of JavaScript enables a variety of techniques for concealing itself from analysis and detection. JavaScript obfuscation can transform a known malicious sample into an undetectable one. Webpack enables bundling benign and malicious scripts and wrangling them to make static analysis harder. The other side of obfuscation is evasions; in addition to making code comprehension (via human or machine) harder, malicious actors have deployed time bombs,

offloading parts of the malicious script to be read from the DOM or via a network request, and dynamic code generation to evade revealing the entire malicious behavior and thus detection. [48]

3 Methodology

This paper provides a methodology for clustering using dynamic features, evaluating how closely the clusters resemble underlying phishing kits, and describing how widespread different adversarial techniques are in the ecosystem. The building blocks of our experiments are browser API execution traces from phishing pages and a ground-truth dataset of pages where we know the underlying phishing kit. In the following section, we describe the experimental setup for gathering this data, the steps we took to aggregate and enrich the execution traces, annotate the clusters based on different techniques, and finally, the steps we took for clustering the data and evaluating them as analogs for phishing kits. A full overview of our crawling and analysis pipeline can be seen in Figure 3.

3.1 Data Gathering

Our crawling infrastructure aims to ingest phishing URLs from upstream providers and output execution traces from the page, as well as a potential kit used for that page.

URL feeds: We gathered phishing pages from a diverse set of phishing feeds, monitoring OpenPhish [47], PhishTank [19], URLScan [54], SMS Gateways [42], PhishDB [40], and APWG [13], based on the availability of the feed and the level of access we had at the time. Every two hours, we checked these feeds for new urls (limited to the last 48 hours) and submitted them into two different crawlers: VisibleV8 and KitPhishr.

VisibleV8: To get execution traces for every script loaded when visiting the page, we used an automated Chromium-based crawler with VisibleV8 patches applied [29]. The browser is being automated to visit the page and take screenshots with puppeteer [25], an NPM package by Google to help in UI/UX testing and browser automation. The patched Chromium crawler uses puppeteer-stealth, a set of configurations to help mask the headless Chrome and Puppeteer itself from detection tools [46]. We initiated the crawls from a network designated for research purposes, for which the ISP would register as 'educational' for any IP intelligence API. We used Catapult [24], a man-in-the-middle proxy, to capture the entire HTTP archive for replayability. The crawler stays on the page for 45 seconds before taking a screenshot to allow scripts to load and start executing; this is consistent with prior work [22, 29].

KitPhishr: While not guaranteed, some malicious actors leave the zip files of the kits used in a discoverable folder on the same server that hosts the website (for example, the Apache document root). KitPhishr [2] is a Go-based URL fuzzer that attempts to identify any leftover zip files on the server. Prior work [36, 44] establishes KitPhishr as a method for collecting and analyzing phishing kits. If successful, it will download the zip file and make a note of the domain from which KitPhishr acquired it.

3.2 Deduplication

Once we have collected browser API traces and potential phishing kits, we post-process the traces into a set of APIs executed in a 1st-party context per page and deduplicate the phishing kits based on archive file similarity.

3.2.1 Trace postprocessing

VisibleV8 has a default log-postprocessor set. These programs take the raw lags generated by the patched Chromium browser and convert it into an organized database, identifying duplicate scripts via SHA3 hash, clearly marking the origin of each script that loaded and isolating which JavaScript API calls are browser API calls defined in the WebIDL file[1] generated while building the patched Chromium. For our analysis, we use a tuple of the original page's URL, the script's URL, and an unordered set of APIs executed by this script.

To not introduce artifacts into our clusters from cloaked pages and 3rd-party scripts like Google Analytics, known to be present in phishing pages [34], we isolate API sets executed by 1st-party scripts (from now on called 1st-party API sets). We establish the root domain as the domain submitted to the feeds and the origin as the domain from which the script is loaded. We consider a script 1st-party only if it is loaded from the domain we acquired from our feeds (root domain). The only exception is when we identify which pages embed a Cloudflare Turnstile script. As some of the Turnstile scripts can be hosted on 'Cloudflare.com.'

Deduplicating kit files: We crawl the URLs with KitPhishr to establish a ground truth dataset with URLs originating from the same kit[2]. For zip files extracted via KitPhishr that have password-based encryption enabled, we use the SHA256 hash of the zip files to identify which domains yielded the same kit. For the remaining zip files, we further de-duplicate them into Kit-Families by looking at them as a set of SHA256 of source files[3]. If the Jaccard index-based similarity ($JI(A,B) = |A \cap B|/|A \cup B|$) of these sets is equal to or greater than 95%, we consider the two kits to be from the same family, and thus group all domains from both kits into the same group.

Adjusting for Cloudflare: We find a high number of anomalies originate from Cloudflare scripts on the same domain as the phishing pages. Since Cloudflare scripts load from a URL with a **'cdn-cgi'** in the URL, for most of our analysis, we

---

[1] https://developer.mozilla.org/en-US/docs/Glossary/WebIDL

[2] For any domain that yeilded two or more zipfiles, we disrecard all

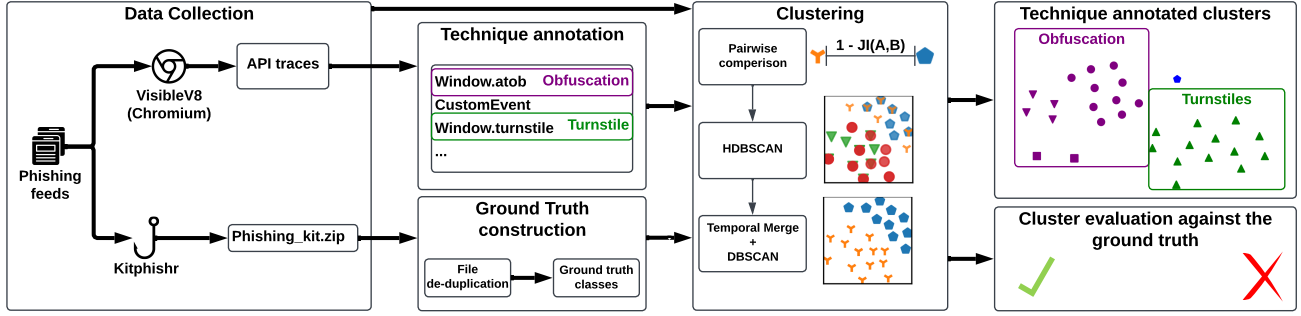[3] Source files identified via a python libmagic module [5]

Figure 3: Crawling and clustering infrastructure

do not consider scripts loaded from that endpoint. However, we discuss the behaviors we could see from these scripts in Section 5.

3.3 Data enrichment

In addition to data from our phishing feeds, we compile API usage from brand's original pages in order to discuss the differences between phishing pages and the login pages they are targeting in Section 6 and we present a mapping of different adversarial technique employed by phishing pages mapped to different API usages.

Technique to Browser API mapping: Client-side JavaScript can engage in data harvesting (exfiltrating information dynamically and not through form submission), evasion (conditional dynamic behavior aimed at hiding functionality or contents), obfuscation (unconditional behavior meant to hinder static analysis), and mimicking (dynamic behavior to make the page more believable, for example, false loading pages, stage by stage data extraction). Based on prior work by Su *et al.* and Zhang *et al.* and manually identifying APIs from the Mozilla Developers Network (MDN) documentation, we present a table mapping standard phishing techniques to browser APIs in Table 1.We leverage the presence of these APIs in the execution traces as a signal of the technique being present in the page. If the page falls within a certain cluster, we mark the entire cluster as using that technique. This is done because we observe phishing pages load scripts that engage in non-deterministic behavior, as well as the difference in the load on our infrastructure, can cause different amounts of work (lines of code) to be executed between different visits. For Cloudflare turnstiles embedding, we use a non-browser API as our detection metric, as Cloudflare's native turnstile script will read the value of `Window.Turnstiles`. For Client-side IP checks, we manually looked through every `Window.fetch` and `XMLHttpRequest.open` argument URL given that argument was present in more then 50 phishing pages, and identified 15 that were IP reputation APIs.

Brand's Original page: OpenPhish and APWG's eCrime Exchange report the brand that a phishing page targets. We selected 5 out of the top 52 brands targeted based on the popularity by page number, as well as brands that represented seasonality targeted sectors (like the IRS or banks) and collected VisibleV8 logs for their home pages and, when applicable, their login pages, to assess how similar phishing pages are to their target pages.

3.4 Identifying kits

We hypothesize that similarity in browser API execution means that the pages originate from the same phishing kit. To test this, we ingest API sets from pages we were able to identify phishing kits for and output a potential clustering of those pages, checking how well the clustering maps back to the ground truth information.

To establish this similarity, we use the Jaccard index on the API sets that 1st-party scripts execute. We use Hierarchical Density-Based Spatial Clustering of Applications with Noise [15] (HDBSCAN), with a minimum cluster size of 2, to cluster the pages with the Jaccard distance as our distance kernel. With HDBSCAN requiring minimal fine-tuning out of the box, and requiring no prior knowledge of the number of clusters we need to look for, we use the ground truth labels from KitPhishr to evaluate the clustering. When ground truth is available, we evaluate the clustering using the **Fowlkes-Mallows Index** [23] and **V-measure**. **V-measure** is a harmonic mean between completeness (all members in a cluster are from the right class) and homogeneity (all members in a cluster are from the same class). V-measure allows tuning a ratio β that prioritizes the score towards one vs the other. More importantly, completeness and homogeneity scores describe how the algorithm is miss-classifying the datapoints [51]. We use sklearn's measure module to calculate all cluster evaluation metrics. [49]. For the much larger set of pages for which we do not have kit labels, we use **silhouette score**[4], a metric

---

[4]While silhouette score is biased against non-convex clusters, based on our results, we do not see it necessary to switch to a density-specific cluster

Table 1: Manual mapping of phishing techniques to browser APIs

Technique	Category	Identifying markers
Fingerprinting extraction	Credential Harvesting	10 Fingerprinting API calls and an exfiltration related API call from [56]
Client-side IP check	Evasion	`Window.fetch` `XMLHttpRequest.open`
Timing bot detection	Evasion	`Performance.now + Timeout`
Encryption	Obfuscation	`Subtlecrypto.decypt`
Encoding	Obfuscation	`TextDecoder.decode` `window.atob`
Dynamic script Evaluation	Obfuscation	`eval`

Basic fingerprinting	Evasion	`HTMLDocument.cookie`
		`HTMLDocument.referrer`
		`Navigator.userAgent`

Dynamic script creation	Evasion	`HTMLScriptElement.text` `HTMLScriptElement.innerHTML`
Cloudflare Turnstiles	Evasion	`Window.turnstile`

Pop-ups	Evasion	`Navigator.requestMIDIAccess`
		`Clipboard.readText`
		`Geolocation.getCurrentPosition`
		`MediaDevices.getDisplayMedia`
		`HID.requestDevice`
		`Window.{confirm|alert|prompt}`
		`Accelerometer|Gyroscope`
		`Window.showModalDialog`
		`MediaDevices.getUserMedia`
		`SyncManager.register`
		`Clipboard.read`
		`Serial.requestPort`
		`USB.requestDevice`
		`Window.queryLocalFonts`
		`Notification.requestPermission`

for how well packed and separated the clusters are, to evaluate the clustering when we do not have ground truth.

To process 533,514 pages in one go, we divide the pages into small time windows and first cluster them into smaller local clusters. Figure 3 shows the breakdown of our methodology. But we first divide our data using a 3-week rolling window; we roll by 2 weeks. However, pages now may exist across two clusters, so once we cluster these chunks using HDBSCAN, with no hyperparameter tuning and a minimum cluster size of 2, we merge any 2 clusters with at least one page in common across the sliding window. We refer to these clusters as 'local clusters'. However, these clusters can not represent phenomena like the re-emergence of clusters, if it happens after 2 weeks. To resolve this, we use the set intersection of APIs (representative API set for the clusters) from every page within a local cluster. We use DBSCAN with a conservative $\epsilon = 0.05$ to merge them. Intuitively, this ϵ represents the maximum dissimilarity tolerated between clusters. We selected this ϵ value by experimenting on the clusters of

---
metric

pages from the ground truth set. When it comes to merging local clusters, $\epsilon = 0.04$ yields fewer clusters and a better silhouette score; however, to avoid data-peeking, as we use the silhouette score to describe the shape of the final clusters, we chose not to determine the value based on the final clusters. We present the clustering metrics from both local clusters and final clusters in Section 5. To avoid calculating a distance matrix on the over four hundred thousand pages successfully clustered, we use the representative set from each local cluster to compute the silhouette score on the final clusters. We separate local clusters labeled as noise in the 2nd cluster step (with DBSCAN and $\epsilon = 0.05$) into single clusters.

Cluster lifetime and deployment diversity: Throughout this work, we refer to cluster lifetime as the time range between when the first page belonging to the cluster is observed on the feed and when the last page belonging to the cluster is observed on the feed. We pull and crawl URLs from the phishing feeds every two hours, meaning that this is an approximation of when the URLs appear on the feeds, with an error of two hours. We measure deployment diversity of the phishing pages by looking at the effective 2nd-level domain (e2LD)

for the URLs. Using the e2LD instead of the entire hostname ensures that pages deployed on 'pages.dev' or 'blogger.com' are considered a single deployment form.

4 Data characterization

4.1 Phishing Kits and Ground truth

We collected 7,273 archive files from XXXX number of domains. Only two were encrypted, which we treated as unique kits based on their SHA256 hash. We further filtered the archive files bu ensuring that they container at least a single code file as identified by Magika. This further reduced our archive count to XXXX. For the final step of filtering, we removed any kit to domain mappings if a domain mapped to multiple archives to remove confusion when multiple kits were deployed on the same domain. From these XXXX we group these into 521 based on SHA256 of code files. We noted that most kits, had a rather extreme

5 Results

In this section, we evaluate our clustering approach against the ground truth, inspect the distribution of ground truth pages in the final clusters, and report temporal patterns and the commonality of phishing behaviors across those clusters. In total, we crawled for *545* days collecting browser traces from 1,328,917 pages, out of which only 533,514 qualified for clustering by executing at least 8 JavaScript APIs in a first-party context. **376,762** did not execute any JavaScript in a first party context and the rest **418,641** did not execute at least 8 distinct browser APIs. Furthermore, 99,464 pages were clustered as noise by HDBSCAN or formed a cluster where all the pages had no common APIs. As shown in Figure 1, in contrast to the number of pages we observe monthly, we can substantially reduce the scope of the phishing ecosystem based on common behaviors. On average, we reduce the monthly traffic to 6% ($\sigma = 2$) of the traffic of urls. Before the 2nd DBSCAN merge of the 26,936 local clusters, Figure 4 shows the distribution of the silhouette score of the local clusters. With an average score of 0.8, these clusters are incredibly well formed and, on average, pages inside a cluster contain 64 APIs in common. Once clustered together, the silhouette score of the final clusters was **0.6**. The average cluster in our dataset contained 49 pages ($\sigma = 468$) and existed for 41 days ($\sigma = 77$).

We collected 2,590 files with KitPhishr for the ground truth dataset, and de-duplicated them into 2,516 phishing kits. Only **519** have at least two URLs that we crawled, for a total of 4,438 pages. When evaluating the ground truth, we dropped the API requirement to four in a first-party context.
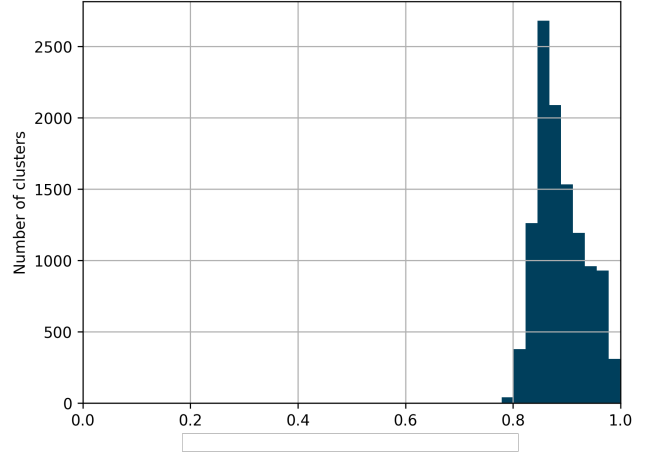

Figure 4: Histogram of the silhouette score of local clusters. This score is calculated including the noise cluster, pages of which are not considered for the merge algorithm

5.1 Characteristics of clusters

Finding 1: *Pages from different kits have an average similarity of 8.5% ($\sigma = 13.7$)*. Figure 5 shows the distribution of Jaccard index-based similarity between pages from the same kit versus different kits. We observe that pages from different kits, even ones with similar themes, rarely exceed 60% similarity, while most pages with the same kit have around 90-100% similarity. We leverage this in our clustering by treating $1 - JI(A, B)$ as a distance kernel of HDBSCAN.

Finding 2: *Browser API usage identify phishing kits apart from each-other*. Clustering pages from 4,448 pages across 521 kits, yield 349 clusters. Evaluating these clusters against the ground truth labels for each page, we find that our clusters have an FMI-based accuracy of *0.93*. The clusters have a V-measure of *0.86*, with compleeteness score of 0.9 and homogeneity score of 0.82.

Pages with ground truth labels for originating kits remain appropriately sorted out in the final clusters. We maintain an FMI of **0.95** and a V-measure of **0.89**, respectively. After manually inspecting a sample of clusters, we observed that these clusters have unique pages across deployment types (AWS, Cloudflare, DigitalOcean, etc.) and languages. For example, **Cluster-e325887b** comprises **487** pages across five unique e2TLDs and contains pages engaging in voice-based phishing attacks (tech support scams) across Japanese, English, and German, varying phone numbers and error messages in each, shown in Figure 9. With over 400 APIs in common, it is clear that these pages' usage of keyboard intercepting APIs, Audio APIs, and Network APIs calls (to **ipwho.is**) for IP intelligence caused the cluster to be formed.

To show how the distance based on the Jaccard index seperated different pages in the larger clusters, Figure 6 shows two different clusters, with example screenshots pulled from both.
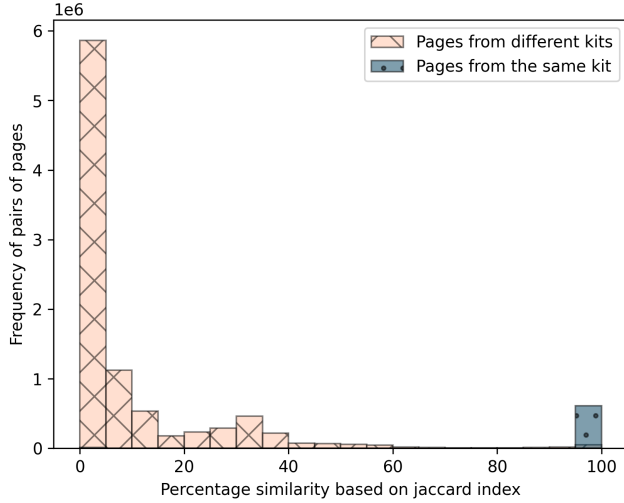
7

Figure 5: Histogram of the Jaccard index-based similarity between any pair of pages belonging to the same kit, versus from different kits

One cluster is from a deployment we can tie to a public GitHub repository (discussed in Finding-11), and the other is a collection of crypto-wallet targeting pages. The figure shows a t-distributed stochastic neighbor (t-SDN) embedding for the distance matrix between all the pages from the two clusters to illustrate better how the pages are separated.

Finding 3: *DOM APIs and Property reads aid in differentiating different pages based on the underlying kit.* With browser APIs, feature reduction becomes an obvious goal. However, removing DOM-related APIs or property reads out of consideration drastically reduces the number of pages we can consider for ground truth evaluation without increasing our overall accuracy. Clustering evaluation metrics for the ground truth dataset with DOM, SVG, and CSS APIs removed results in an FMI of 0.93 and a V-measure of 0.86. When all property reads were removed, we saw an FMI and V-measure of 0.93 and 0.85, respectively. In both cases, we can cluster fewer pages and thus identify fewer kits, but we do not see a significant improvement in clustering accuracy.

Finding 4: *69% of clusters contain urls only marked by a single target brand by our threat intel sources*[5]. This phenomena with the ground truth dataset, as URLs for 90% of the kits were targetting a singular brand. Together, this provides strong circumstantial evidence that deployed kits are increasingly becoming brand-specific. 1,253 clusters (19%) of the clusters had two brand labels. However, the most popular combination of these was "Meta/Facebook", "Facebook/Instagram", "National Police Agency JAPAN/Facebook", and "holiganbet/jojobet", keeping the parent organization of the target the same in the majority of the cases. Manual examination of clus-

---
[5]We did not include clusters in this count that had no brand-labeled urls in them

ters with "National Police Agency JAPAN/Facebook" brand labels revealed shopping pages in Japanese to be marked with that label incorrectly from our data feeds. The cluster with the most diverse set of brand labels had 14 unique brand labels, which was a cluster with 12,467 pages with simple sign-in pages that exfiltrated information using client-side registered event listeners. Furthermore, we find that 16,100 pages (3% of the pages observed spanning 313 clusters) come from phishing kits collected using KitPhisher; however, we could not pull the kit from the URL in all pages. Figure 7 shows a histogram of the number of unique brand labels observed per cluster.

5.2 Temporal patterns of clusters

Finding 5: *Top 50 clusters by page count, account for over 40% of the phishing traffic observed. 40.7%* pages are sorted into one of the 50 top clusters. The 10 clusters by page count are 20.8% of the total pages alone. We provide a breakdown of these clusters along with manual labeling of what campaigns they correspond to in Figure 8. We note that one of the clusters, which we labeled dynamically generated, turned out to be a noisy cluster of simple pages that dynamically generate the page using JavaScript, without any sophisticated client-side behavior, based on their browser API traces. The E-commerce cluster shown in Figure 8 has a significant seasonality, with an auto-correlation function showing significance at lags 7,14, and 21; meaning the appearance of the clusters on the feeds happens every week. However, further investigation showed that the majority of seasonal clusters are similar e-commerce phishing clusters, with vastly different dynamic behavior, allowing us to conclude that the seasonality in the majority of the clusters is due to regular reporting by threat-intelligence sources to our feeds, and not seasonal deployments of kits.

Finding 6: *72% of the clusters (144,695 pages), are only seen for a single month by phishing feeds.* Lack of re-emergence for these clusters signals about the existence of throw-away kits, or the possibility that these clusters engage in profiling of phishing feed detectors as described by [12]. 15% of the clusters (14,555 pages) we observe are only seen for one day, along with brandless bank pages and a phishing page impersonating the government of Korea. Meanwhile, 1,282 clusters (11%) have lifetimes longer than 100 days. Some clusters with a lifetime greater than 100 days still only deploy a few pages (less than 1 page every 10 days). We use this as a heuristic to identify 251 clusters 're-emergence' through our observation period. It should be noted that every day, there is a new cluster that appear on the phishing feeds based on behaviors, however, from Figure 1, the total number of clusters active on all of our phishing feeds is much more manageable, especially when considered that the clusters can be queried based on their dynamic behaviors.
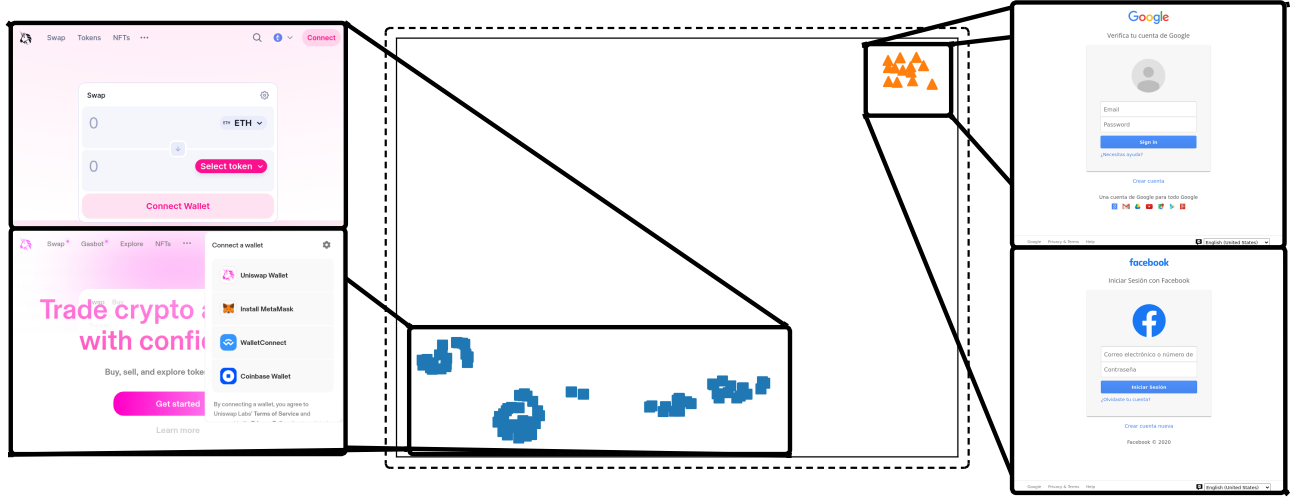
Figure 6: Example of pages from two different pages and embedding between them. We used t-distributed stochastic neighbor embedding to visualize the distance matrix between all pages in a 2-D plane. The clusters presented here are from Ethereum wallet pages with a few variations of the landing page (▲) and pages that descend from an open-source phishing kit discussed in Section 5 (■). We should note that while the pages have different logos on the right-hand side, the bottom left corner still reads "Google" on those pages.
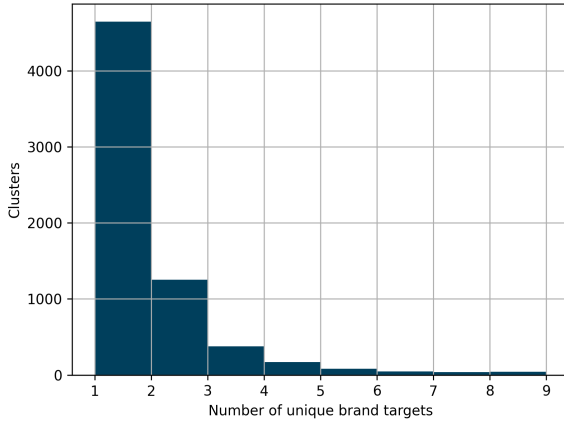

Figure 7: Distribution of unique brand labels per cluster (only counting clusters that had at least one page with a brand label)

5.3 Phishing Techniques across clusters

Finding 7: *UI interactivity and fingerprinting are a near-universal behavior across clusters*. Multi-stage phishing pages are very well documented in prior work, and we find that the majority of clusters (91%) register a click event listener using JavaScript. Though this could be as simple as submitting credentials using JavaScript, this highlights the need for researchers to augment their crawlers in the future to extract better and more complete execution traces from websites. We

Table 2: Breakdown of the obfuscation techniques observed in our dataset

Obfuscation techniques	Pages	Clusters
Window.atob	61,125	1,455
eval	14,561	982
Textdecoder.decode	11,113	534
SubtleCrypto.decrypt	1,185	36

split fingerprinting into two categories, basic and advanced. Basic fingerprinting, which follows the list of APIs identified by Zhang *et al.* in [62] was present in 80% of the clusters (over 300,000 pages), and Advance fingerprinting (measured by at least 5 APIs idenfied by Su *et al.* in [56]) show up in 70% of the clusters. Together, 85% of clusters (9,572 clusters, 313,212 pages) exhibit some form of fingerprinting.

Finding 8: *Fingerprint exfiltration, obfuscation, and bot detection are widespread across clusters*. While fingerprinting is near universal, we find that a smaller fraction of the clusters employ obfuscation, fingerprint exfiltration, and timing for bot detection. Prior work has shown interest in these behaviors [35, 48, 52, 63], meaning kits that forgo this may be rudimentary either by negligence or design, to avoid static and dynamic phishing detection based on JavaScript features. Dropping anti-bot detection features has been observed before, with an Office365 phishing kit (dubbed Tycoon2FA), opting to remove CloudFlare Turnstile integration, as it was being used as a feature for detection [8].

```
await this.$http({
  method: "get",
  url: "https://api.ipregistry.co/?key=" +
  ↪ this.key
}).then(e => {
  const s = e.connection.type,
      c = ["cdn", "hosting", "education"];
    /*omitted for brevity*/
    if (c.indexOf(s) != -1)
        /*Redirect*/
})
```

Listing 1: Example of IP-based cloaking by using a 3rd-party reputation API. The following example specifically cloaks away from educational networks like ours.

Another common tactic for bot detection is timing-based checks; this can be done via Browser APIs by calling `Performance.now` right before and inside of a `Window.setTimeout` statement to measure the time differential between setting the timeout and its triggering. We find that 22% of clusters call `Performance.now` in conjunction with `setTimeout`.

On the other hand, 31% of clusters (2,395) employ some form of obfuscation. We present the breakdown of all obfuscation techniques in Table 2 and as we can see, eval and Base64 encoding were the most popular ways of obfuscation. Despite the best recommendations to web developers to avoid using 'eval' [9], JavaScript's eval function remains a favorite for obfuscation and evasions [48]. Sometimes, a script is executed via 'eval()', which evaluates yet another script itself; we measure this phenomenon as a level in **eval-depth**. We find that 48 clusters have pages that go to eval-depth 3. However, this seems to be a side-effect of embedding the phishing pages (mainly ones targeting Facebook) in Blogger.com pages.

Finding 9: *While rare, client-side IP reputation checks are present across multiple clusters*. While only present in 504 clusters (19,869 pages), we identify 15 unique IP reputation APIs used by phishing pages as soon as the page loads. We present a full breakdown in Table 3. While not the most popular, *api.ipregistry.co* presents an interesting case study, as it enables the identification of educational networks. Manual examination of pages from these clusters reveals snippets similar to Listing 1, which conditionally chooses to redirect away from cloud hosting providers, content delivery networks, and educational networks, like the one we performed the crawls through. However, due to them employing other browser APIs, before the cloaking behavior, we are still able to cluster the pages based on the initial logic of the landing page.

Finding 10: *Pop-UP APIs are declining in usage*. We see only 104 clusters (1,323 pages) call out to pop-up requesting APIs. Among these, Geolocation.getCurrentPosition (55 clusters) was the most popular. While requiring a pop-up to interact

Table 3: List of all API endpoints that client-side code reaches out for IP intelligence.

API url	clusters	pages
api.db-ip.com	40	3,124
api.geoapify.com	3	63
api.ipapi.com	5	143
api.ipgeolocation.io	21	96
api.ipify.org	177	7,417
api.ipregistry.co	9	3,995
freeipapi.com	38	6,444
geolocation-db.com	14	492
geolocation.onetrust.com	37	364
get.geojs.io	14	753
ipapi.co	106	735
ipinfo.io	65	1,963
ipwho.is	47	798
pro.ip-api.com	20	101

with, this API can also play a crucial role in cloaking, as any VPN or proxy does not mask the results.

We observe a smaller fraction of the ecosystem (16 clusters, 148 pages) than [62] employs this cloaking technique, especially when it comes to triggering a notification pop-up to verify user interaction. This could be a result of Firefox, citing low engagement with the notifications, started requiring user interaction to trigger the popup [41] at the end of November 2019, when crawlphish's data collection ended. Chrome has since discussed modifying the notification API to make the request less disruptive to the user experience [3]. The lack of pop-up requests could also be explained by our **Finding 9** regarding usage of IP intelligence APIs or by the overwhelming amount of the pages (67%) registering at least one HTMLElement event handler, which would be classified as a *Click-through* by Crawlphish's taxonomy. We report a full breakdown of APIs related to the Crawlphish categorization of client-side cloaking in Table 5.

Finding 11: *Mouse Detection API calls and Cloudflare Turnstile embedding are specific to a small group of clusters*. While supported by most modern browsers, we see a very rare use of Mouse Detection APIs. Only 35 clusters employ mouse detection-related APIs. Two of these clusters[6] are from an open-source phishing kit, leveraging botguard, and are from a public GitHub repository, which was last updated in 2017[7]. We see these clusters deploy across 17 unique domains, starting from 2023-10-07 all the way to 2024-07-19. **7 clusters (181 pages)** embed a Cloudflare turnstile check in their page; some domains are not hosted on Cloudflare. It should be noted that in the case of redirection. At the same time, we discussed the presence of WebAssembly-based captchas for bot detection,

---

[6]Split due to infrastructure inability causing crashes in early crawls
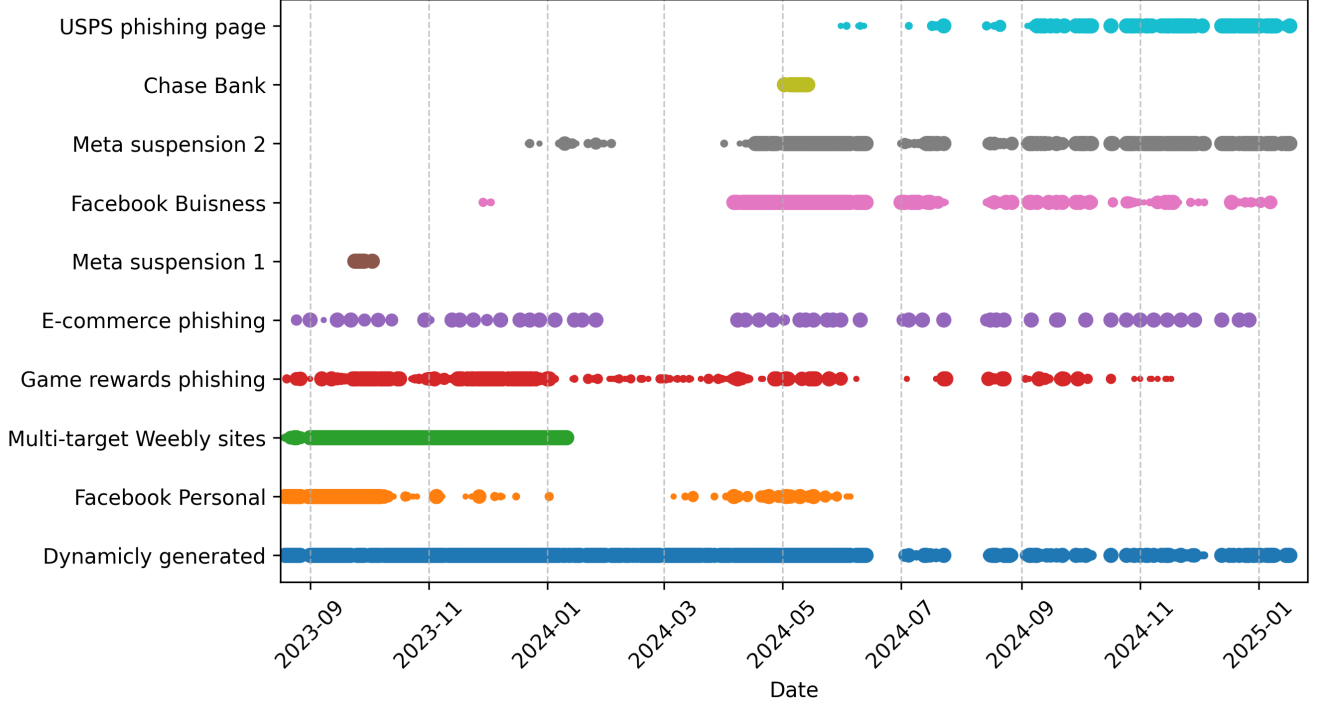[7]https://github.com/ashanahw/Gmail_Phishing

Figure 8: Timeline of the top 10 clusters based on the number of pages

embedding a Cloudflare Turnstile check allows the phishing kit authors to offload bot detection to a well-established ecosystem. Recently, analyses have identified high-value phishing kits with this behavior [4]. However, subsequent analysis of the same threat actor identified a shift from turnstiles to HTML-Canvas drawn captchas.

Finding 12: *The phishing ecosystem consists of clusters that utilize both cutting-edge, experimental browser APIs, and extremely deprecated APIs.* We find 421 clusters (8,270 pages) that utilize *NavigatorUAData.getHighEntropyValues* for fingerprinting and *Keyboard.lock* to restrict user input, APIs not fully supported by Firefox and Safari. We identified 10 clusters across 4,433 pages that used *Scheduling.isInputPending*, however, upon closer inspection, these were not pages using a novel kit, but rather pages that used Google Sheets to construct their landing page. On the other hand, 25% of the clusters spanning (47,001) pages use a deprecated API.

While not experimental, WebAssembly is still a relatively modern web practice. We find a total of 199 clusters (5,107 pages) that use WebAssembly related APIs. Upon manual inspection of 33 unique WASM modules on these pages, we identify bot-detection, in most cases, by using FriendlyCaptcha, as the most common use case for WebAssembly in phishing.

Figure 10 shows the confusion matrix between the techniques we enumerated and the size of the clusters. We see no noticeable difference in techniques regarding cluster size, except for a higher percentage of large clusters using client-side

IP checks. We also observe that pages that employ experimental APIs also tend to include a deprecated API call, which aligns with their usage for browser fingerprinting, rather than novel cloaking logic.

6 Discussion

6.1 Why clusters?

As shown in the examples in this paper, phishing feeds can be a noisy data source for studying the overall ecosystem. From e-commerce pages to mass-spammed USPS and EZ-parking phishing pages, differentiating between a handful of pages of varying sizes becomes crucial for efficiency. This paper demonstrates that browser API usage can closely proxy the phishing kit (providing commonality for the page's origin) at best and group via standard client-side techniques at worst.

Furthermore, cluster membership with a page of a known kit and a high number of common cluster APIs enable the inferral of server-side behavior of the page by propogating threat intelligence from one page to all in the same cluster. As established by Oest *et al.*, phishing kits thrive on including up-to-date IP blocklists. However, this blocklist is not always included in the client-side code. Looking at the pages as is, without any forced execution, allows for easier triaging of failed detections and crawls as it clusters them all together.

Table 4: Comparison of browser APIs executed by phishing pages against the original login page of the brand they are targetting

Brand	Phishing pages observed	Avg % Similar	Std Dev (%)	Perfect Subset	50%+ Match
Facebook	338,686	11.6	9.6	0	24,039
USPS	66,692	10.3	10.5	0	53
Meta	51,969	12.4	7	0	2,605
Microsoft	3,768	11.8	8.7	0	419
IRS	1,207	11.5	9	0	61

6.2 What makes a kit identifiable?

Sets of browser APIs used are not the most granular method to describe pages; however the more sophisticated the phishing kit becomes, the easier it is to spot by just the browser APIs it uses, as everything from how it chooses to evade researchers to how the UI libraries they use build the DOM, can make it stand out. In plain JavaScript, without browser APIs, a page can not reach out to a C2 server, gather browser fingerprints, dynamically generate or cloak page contents, request browser pop-ups, or even respond to a button click. In the case of the pages clustered in Figure 9, the APIs *Keyboard.lock*, *HTML-Document.onkeydown* for keyboard locking, *Window.atob* for obfuscation, and a handful of fingerprint APIs and DOM APIs for dynamic content generation, set these pages apart from the other clusters.

Finding 13: *Phishing pages vary wildly from the brand that they are mimicking*. We collect browser API traces from Facebook, USPS, Meta, Microsoft, and IRS login pages and compare them to their phishing counterparts. The average similarity of the APIs executed by the phishing page and the original page is 11%, indicating that browser APIs do not relate to the target page. We found no pages where the original page's API set was a subset of the phishing page's API set, and in 5% of the cases, the phishing page executed at least half of the APIs from the original page. We report per-brand findings in Table 4 and note that the least similar brand was USPS, which could be the result of USPS phishing pages being multistage pages requiring user interaction and targeting credit card information [17].

This, however, does not indicate that browser API usage alone is a good indicator of maliciousness. All of the techniques described in Section 3 are common throughout the web. Phishing pages are singular in their purpose, and their choice to engage in these techniques sets them apart from one another.

6.3 Dynamic analysis combats evasions

While dynamic analysis suffers from drawbacks we discuss in Section 8, it allows us to sidestep a subset of sophisticated obfuscation techniques. Even something as nuanced as an AES-encrypted script are logged by VisibleV8 as long as the behavior is triggered and the script calls out to browser APIs. While not incorporated by our work, prior work has shown

forced execution [48, 62] to ensure that dynamic analysis effectively extracts all behaviors.

7 Related Work

7.1 Phishing

Phishing detection: There is a wealth of research on phishing detection as the tug-of-war between adversaries and security professionals continues. Recently, Liu *et al.* and Abdelnabi *et al.* have deployed vision-based techniques to detect phishing pages [11, 36, 37]. [37] also presents over 6,000 phishing kits analyzed as part of the work. With adversarial attacks ensuring that the page looks different to crawlers and analysis, some have turned to extracting features from the URLs themselves, more recently via LLMs in [18, 30] and earlier via statistical models and machine learning in [32, 50, 55, 61]

Studying and combating adversarial techniques: Divakaran *et al.* in [21] reaffirms the need to keep up with the latest adversarial techniques to build better detection systems for phishing. Prominent work in this area includes [62] by Zhang *et al.*, which uncovered and categorized many novel client-side techniques by forcing the execution of phishing pages to trigger the cloaking behavior. Acharya *et al.* in [12] uncovered that phishing pages can successfully evade blocklists by knowing how to identify their crawlers, and Oest *et al.* in [43] demonstrated that cloaking from non-mobile based devices as a phishing page can ensure that attacker's page goes unmarked by the blocklists for more than 48 hours.

Kondracki *et al.* in [31] uncovered a massive blindspot of the phishing detection ecosystem that was Man-in-the-middle phishing kits. Kits that would transparently forward the victim's connection to the target page, mimicking brand logos on pages like Outlook without any configuration. Fortunately, the authors addressed the blind spot by demonstrating that these proxies remain fingerprintable using TLS fingerprinting. [60] proposed a similar attack, however, one that used JavaScript and NoVNC to trick the user into signing into their account through a VNC session in their browser.

With adversaries becoming creative with their evasions and obfuscation techniques, some novel defenses have also opted to think outside the box. Zhang *et al.* in [64] proposed a phishing defense solution that leverages the high likelihood of a phishing page cloaking away from a crawler to the de-

fender's advantage. They demonstrated that a web browser configured to look like a crawler triggers a cloaking response from phishing pages, ensuring that victims never see the page while maintaining compatibility with all of Alexa's Top One Million websites. Meanwhile, using CAPTCHAs [59] utilized vision-based models to combat phishing pages.

To better understand why, other than cloaking, phishing pages may choose to fingerprint, Lin *et al.* in [35] showed that browser fingerprints could be successfully used to bypass multi-factor authentication, a system meant to be a last line of defense against stolen credentials, for 10 out of 16 websites that provide popular services.

Phishing kits: Much can be studied about the phishing ecosystem via phishing kits. Cova *et al.* in [20] uncovered that most "free" phishing kits contain a backdoor, effectively serving as a way to offload the deployment of a campaign to a 3rd party while siphoning off their stolen credentials.

Similar to our goals, PhishKitA [16] uses a dataset of phishing kits gathered through **KitPhisher** and a collection of features extracted from the HTML DOM to classify websites into their matching kit. They achieved an F1 score 0.91 when classifying 2,000 pages (1,141 benign and 859 phishing) from features extracted from their kit. However, their multi-class classifier for identifying the kit only achieved an F1 score of 0.39. Merlo *et al.* in [39] further expanded on our understanding of phishing kit lineages by looking at over 20,000 phishing kits and identifying, via token similarity, most of them as clones of one another or previously encountered kits. Prominent work in extending our understanding of phishing attacks includes Han *et al.* in [26], where they monitor the deployment of phishing kits by adversaries that compromise vulnerable web servers by hosting a well-sandboxed honeypot. They collected 643 phishing kits and established that kits take minutes to install and test and can remain undetected for weeks. Using these kits, they were also able to identify evasion techniques used by these kits, like path randomization per visit, which back then was enough to bypass Google SafeBrowsing.

In [44], Oest *et al.* manually analyzed phishing kits to establish the taxonomy for server-side cloaking, and in [14], Bijmans *et al.*, after collecting phishing kits by watching TLS transparency logs to identify Dutch brank phishing domains, manually created a fingerprint from static features to analyze their prevalence in the wild.

More recently, Lee *et al.* in [33] provides a server-side script (PHP) level analysis of phishing kits, finding that dynamically generated URLs are still standard in the ecosystem and observing seasonality in the kits they were able to obtain.

Extending the understanding of the phishing ecosystem:

Similar to our methods, Rola *et al.* in [52] deployed a modified Chromium browser to gather data and analyze phishing website browser APIs utilizing a pre-selected API list focused on first and third-party scripts for phishing pages. They find that the majority of the most visited phishing pages (identified via browser telemetry data) deploy fingerprinting scripts, sometimes varying from the ones of the original brand they portray. At the same time, they accessed this at a script level, reinforcing our finding that phishing pages vary vastly from their original page.

Oest *et al.* in [45] demonstrates the full lifecycle of a phishing campaign by employing the fact that phishing pages often copy assets from the target domain and refer the victim back to the original page afterward. By collaborating with a significant financial institution, they developed a framework for leveraging this data to track a phishing page from its deployment to blocklists flagging the page as phishing. [45] observed all techniques highlighted by prior work: cloaking, user-specific URL generation, man-in-the-middle proxies, and short-lived bursty attacks. Expanding our understanding of the victim experience on a phishing website, Subramani *et al.* in [57] developed a crawler.

7.2 Dynamic analysis of webpages

Our work shares a methodology for dynamic analysis enabled by web measurement frameworks like OpenWPM [22] and VisibleV8 [29]. Su *et al.* used VisibleV8 traces and taint analysis in [56] to discover emerging fingerprinting techniques. Sarker *et al.* used VisibleV8 to create an oracle for detecting obfuscation [53]. Such an oracle was made possible by the observation that VisibleV8 marks the execution of an API at a given source line. At the same time, obfuscation techniques ensure that the API is not textually available there. And Pantelaios *et al.* used a combination of VisibleV8 and force execution modifications to the Chromium engine to identify and defeat JavaScript evasion techniques while also leveraging API traces and clustering to identify previously unlabeled malicious extensions [48]. Iqbal *et al.* used OpenWPM to capture execution traces from tranco top 100K URLs of Tranco, training a classifier on a mixture of dynamic and static features extracted from the JavaScript's AST and execution traces, respectively, to achieve a 99.8% accuracy in identifying fingerprinting scripts online.

Our work differs from prior work in multiple ways. To date, we are the most successful in identifying phishing kits via the page's dynamic features; moreover, our differentiation is entirely automated, requiring no prior rule-based identification of kits. While we integrate many of the techniques annotated from prior work, we contribute an up-to-date understanding of their distribution in their ecosystem, and we do so at the cluster level, which we argue is more likely to control for multiple deployments of the same sophisticated kit.

8 Limitations and Future Work

8.1 Variance

We collect API traces from these pages by remaining on these pages for 45 seconds. While this aligns with what prior work employs, the lack of force execution, page interactions, and different system loads between runs introduces noise into our measurement. As highlighted by other work [57], modern phishing pages include a multi-step user experience, and future work can aim to extrapolate the techniques used by pages at different stages via instrumented crawlers that interact with the pages.

8.2 Bias sampling

While most phishing pages have been shown to race the clock against being listed on a blocklist eventually, prior work has shown that these blocklists can have blind spots. While we follow best practices in diversifying our sources for phishing pages, we remain bound to the portion of the ecosystem that our orecles have eventually discovered.

8.3 Limited Kit sophistication

The phishing kits in this paper were collected using **KitPhishr**; while this has been done in prior research, this requires a misconfiguration on the side of the phisher, thus the bias towards PHP, as it is relatively easy to misconfigure a PHP deployment to allow path traversal. However, modern backend web technologies, like ExpressJS, Flask, and even Go's built-in HTTP server, would nullify this weakness. The dataset of diverse phishing kits is hard to come by. It could only be obtained through gaining trust in communities where they are freely shared (i.e., Telegram groups, forums), direct purchasing, or a honeypot setup similar to [26] but one that would enable deployment of any kits.

8.4 Non-clustered pages

For this paper, we ignored any page that did not execute at least eight distinct browser APIs in a first-party context. 376,762 pages did not execute any JavaScript at all, meaning that they are limited to server-side cloaking, and dynamic page generation on the server-side for obfuscating static features, while a possibility, we consider server-side techniques out of scope for our research purposes. Most of the chunk of pages were disqualified for not executing enough browser APIs or executing many of them in a first-party context. Including all party scripts would have muddied the similarities with shared usage of standard libraries, which has been documented in prior work [34], and allowing executions from redirected pages could introduce features into our clustering from the benign pages they are redirecting to. Phishing pages that employ server-side cloaking could redirect the victim to a benign page that loads scripts in a first-party context relative to itself; however, this causes pages to be grouped together based on the redirection target. This can be an avenue for future work to explore, as including first-party scripts relative to the origin (current page loaded) would become more valuable if the crawler is automated to interact with the phishing page, to extract behaviors.

9 Conclusion

In this paper, we provided a workflow for researchers and analysts to automatically differentiate between a collection of phishing pages, based on a common underlying kit or shared techniques, if the behaviors are too generic. We show an accuracy of 97% against a dataset of pages and kits collected from the wild. With a curated mapping of techniques to browser APIs and over 500,000 pages in which we identify *11,377* clusters, we explore what techniques are universal, widespread across kits, or kit-specific.

References

[1] APWG | Phishing Activity Trends Reports. https://apwg.org/trendsreports/.

[2] cybercdh/kitphishr: A tool designed to hunt for phishing kit source code. https://github.com/cybercdh/kitphishr.

[3] Introducing quieter permission UI for notifications. https://blog.chromium.org/2020/01/introducing-quieter-permission-ui-for.html.

[4] Latest Alerts and Advisories | NJCCIC. https://www.cyber.nj.gov/Home/Components/News/News/1424/214.

[5] Python-magic: File type identification using libmagic. http://github.com/ahupp/python-magic.

[6] Sophisticated Spearphishing Campaign Targets Government Organizations, IGOs, and NGOs | CISA. https://www.cisa.gov/news-events/cybersecurity-advisories/aa21-148a.

[7] Threat Actor Leverages Compromised Account of Former Employee to Access State Government Organization | CISA. https://www.cisa.gov/news-events/cybersecurity-advisories/aa24-046a.

[8] Tycoon2FA New Evasion Technique for 2025. https://www.trustwave.com/en-us/resources/blogs/spiderlabs-blog/tycoon2fa-new-evasion-technique-for-2025/?hs_amp=true.

[9] Eval() - JavaScript | MDN. https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/eval, January 2025.

[10] Unit 42. Threat Actor Groups Tracked by Palo Alto Networks Unit 42. https://unit42.paloaltonetworks.com/threat-actor-groups-tracked-by-palo-alto-networks-unit-42/, June 2024.

[11] Sahar Abdelnabi, Katharina Krombholz, and Mario Fritz. VisualPhishNet: Zero-Day Phishing Website Detection by Visual Similarity. In *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security*, CCS '20, pages 1681–1698, New York, NY, USA, November 2020. Association for Computing Machinery.

[12] Bhupendra Acharya and Phani Vadrevu. {PhishPrint}: Evading Phishing Detection Crawlers by Prior Profiling. pages 3775–3792.

[13] Anti-Phishing Working Group (APWG). eCrime Exchange (eCX). https://apwg.org/, 2025.

[14] Hugo Bijmans, Tim Booij, Anneke Schwedersky, Aria Nedgabat, and Rolf van Wegberg. Catching Phishers By Their Bait: Investigating the Dutch Phishing Landscape through Phishing Kit Detection. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 3757–3774, 2021.

[15] Ricardo J. G. B. Campello, Davoud Moulavi, and Joerg Sander. Density-based clustering based on hierarchical density estimates. In Jian Pei, Vincent S. Tseng, Longbing Cao, Hiroshi Motoda, and Guandong Xu, editors, *Advances in Knowledge Discovery and Data Mining*, pages 160–172, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.

[16] Felipe Castaño, Eduardo Fidalgo Fernández, Rocío Alaiz-Rodríguez, and Enrique Alegre. PhiKitA: Phishing Kit Attacks Dataset for Phishing Websites Identification. *IEEE Access*, 11:40779–40789, 2023.

[17] SANS Internet Storm Center. USPS Phishing Scam Targeting iOS Users. https://isc.sans.edu/diary/30078.

[18] Daiki Chiba, Hiroki Nakano, and Takashi Koide. DomainLynx: Leveraging Large Language Models for Enhanced Domain Squatting Detection. *ArXiv*, abs/2410.02095, 2024.

[19] Cisco Talos Intelligence Group (Talos). Phishtank. https://phishtank.org/, 2025.

[20] Marco Cova, Christopher Kruegel, and Giovanni Vigna. There is No Free Phish: An Analysis of "Free" and Live Phishing Kits.

[21] Dinil Mon Divakaran and Adam Oest. Phishing Detection Leveraging Machine Learning and Deep Learning: A Review. *IEEE Security & Privacy*, 20(5):86–95, September 2022.

[22] Steven Englehardt and Arvind Narayanan. Online Tracking: A 1-million-site Measurement and Analysis. In *Proceedings of ACM CCS 2016*, 2016.

[23] Edward B Fowlkes and Colin L Mallows. A Method for Comparing Two Hierarchical Clusterings. *Journal of the American statistical association*, 78(383):553–569, 1983.

[24] Google Inc. Catapult. https://chromium.googlesource.com/catapult/, 2025.

[25] Google Inc. Puppeteer. https://pptr.dev/, 2025.

[26] Xiao Han, Nizar Kheir, and Davide Balzarotti. PhishEye: Live Monitoring of Sandboxed Phishing Kits. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, CCS '16, pages 1402–1413, New York, NY, USA, October 2016. Association for Computing Machinery.

[27] Microsoft Threat Intelligence. New sophisticated email-based attack from NOBELIUM. https://www.microsoft.com/en-us/security/blog/2021/05/27/new-sophisticated-email-based-attack-from-nobelium/.

[28] Microsoft Threat Intelligence. Frankenphish: TodayZoo built from other phishing kits. https://www.microsoft.com/en-us/security/blog/2021/10/21/franken-phish-todayzoo-built-from-other-phishing-kits/, October 2021.

[29] Jordan Jueckstock and Alexandros Kapravelos. VisibleV8: In-browser Monitoring of JavaScript in the Wild. In *Proceedings of the ACM Internet Measurement Conference (IMC)*, 2019.

[30] Koide, Takashi and Fukushi, Naoki and Nakano, Hiroki and Chiba, Daiki. Phishreplicant: A language model-based approach to detect generated squatting domain names. In *Proceedings of the 39th Annual Computer Security Applications Conference*, ACSAC '23, page 1–13, New York, NY, USA, 2023. Association for Computing Machinery.

[31] Brian Kondracki, Babak Amin Azad, Oleksii Starov, and Nick Nikiforakis. Catching Transparent Phish: Analyzing and Detecting MITM Phishing Toolkits. In *Proceedings of the 2021 ACM SIGSAC Conference on Computer*

and Communications Security, CCS '21, pages 36–50, New York, NY, USA, November 2021. Association for Computing Machinery.

[32] Hung Le, Quang Pham, Doyen Sahoo, and Steven C. H. Hoi. URLNet: Learning a URL Representation with Deep Learning for Malicious URL Detection. *ArXiv*, abs/1802.03162, 2018.

[33] Woonghee Lee, Junbeom Hur, and Doowon Kim. Beneath the phishing scripts: A script-level analysis of phishing kits and their impact on real-world phishing websites. In *Proceedings of the 19th ACM Asia Conference on Computer and Communications Security*, pages 856–872. ACM.

[34] Kyungchan Lim, Jaehwan Park, and Doowon Kim. Phishing Vs. Legit: Comparative Analysis of Client-Side Resources of Phishing and Target Brand Websites. In *Proceedings of the ACM Web Conference 2024*, WWW '24, New York, NY, USA, 2024. Association for Computing Machinery.

[35] Xu Lin, Panagiotis Ilia, Saumya Solanki, and Jason Polakis. Phish in Sheep's Clothing: Exploring the Authentication Pitfalls of Browser Fingerprinting. pages 1651–1668.

[36] Ruofan Liu, Yun Lin, Xianglin Yang, Siang Hwee Ng, Dinil Mon Divakaran, and Jin Song Dong. Inferring phishing intention via webpage appearance and dynamics: A deep vision based approach. pages 1633–1650.

[37] Ruofan Liu, Yun Lin, Yifan Zhang, Penn Han Lee, and Jin Song Dong. Knowledge expansion and counterfactual interaction for {Reference-Based} phishing detection. pages 4139–4156.

[38] Heather McCalley, Brad Wardman, and Gary Warner. Analysis of Back-Doored Phishing Kits. In Gilbert Peterson and Sujeet Shenoi, editors, *Advances in Digital Forensics VII*, pages 155–168, Berlin, Heidelberg, 2011. Springer.

[39] Ettore Merlo, Mathieu Margier, Guy-Vincent Jourdan, and Iosif-Viorel Onut. Phishing kits source code similarity distribution: A case study. In *2022 IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pages 983–994. ISSN: 1534-5351.

[40] Mitchell Krog and Nissar Chababy. PhishingDB. https://github.com/Phishing-Database/Phishing.Database, 2025.

[41] Mozilla. Restricting Notification Permission Prompts in Firefox. https://blog.mozilla.org/futurereleases/2019/11/04/restricting-notification-permission-prompts-in-firefox, November 2019.

[42] Aleksandr Nahapetyan, Sathvik Prasad, Kevin Childs, Adam Oest, Yeganeh Ladwig, Alexandros Kapravelos, and Bradley Reaves. On SMS Phishing Tactics and Infrastructure. In *2024 IEEE Symposium on Security and Privacy (SP)*, pages 169–169. IEEE Computer Society, 2024.

[43] Adam Oest, Yeganeh Safaei, Adam Doupe, Gail-Joon Ahn, Brad Wardman, and Kevin Tyers. PhishFarm: A Scalable Framework for Measuring the Effectiveness of Evasion Techniques against Browser Phishing Blacklists. In *2019 IEEE Symposium on Security and Privacy (SP)*, pages 1344–1361, San Francisco, CA, USA, May 2019. IEEE.

[44] Adam Oest, Yeganeh Safei, Adam Doupe, Gail-Joon Ahn, Brad Wardman, and Gary Warner. Inside a Phisher's Mind: Understanding the Anti-Phishing Ecosystem through Phishing Kit Analysis. In *2018 APWG Symposium on Electronic Crime Research (eCrime)*, pages 1–12, San Diego, CA, May 2018. IEEE.

[45] Adam Oest, Penghui Zhang, Brad Wardman, Eric Nunes, Jakub Burgis, Ali Zand, Kurt Thomas, Adam Doupé, and Gail-Joon Ahn. Sunrise to Sunset: Analyzing the End-to-end Life Cycle and Effectiveness of Phishing Attacks at Scale. pages 361–377.

[46] Berstend on Github. Puppeteer stealth plugin. https://www.npmjs.com/package/puppeteer-extra-plugin-stealth, 2025.

[47] OpenPhish. OpenPhish. https://www.openphish.com/, 2025.

[48] Nikolaos Pantelaios and Alexandros Kapravelos. FV8: A Forced Execution JavaScript Engine for Detecting Evasive Techniques. In *Proceedings of the USENIX Security Symposium*, August 2024.

[49] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

[50] Ch. Chakradhara Rao, A.V.T. Raghav Ramana, and B. Sowmya. Detection of phishing websites using hybrid model. 2018.

[51] Andrew Rosenberg and Julia Hirschberg. V-measure: A conditional entropy-based external cluster evaluation measure. In *Conference on Empirical Methods in Natural Language Processing*, 2007.

[52] Iskander Sanchez-Rola, Leyla Bilge, Davide Balzarotti, Armin Buescher, and Petros Efstathopoulos. Rods with Laser Beams: Understanding Browser Fingerprinting on Phishing Pages. pages 4157–4173.

[53] Shaown Sarker, Jordan Jueckstock, and Alexandros Kapravelos. Hiding in Plain Site: Detecting JavaScript Obfuscation through Concealed Browser API Usage. In *Proceedings of the ACM Internet Measurement Conference (IMC)*, page 648–661, October 2020.

[54] SecurityTrails. Urlscan. https://urlscan.io/, 2025.

[55] Hossein Shirazi, Bruhadeshwar Bezawada, and Indrakshi Ray. "kn0w thy doma1n name": Unbiased phishing detection using domain name based features. *Proceedings of the 23nd ACM on Symposium on Access Control Models and Technologies*, 2018.

[56] Junhua Su and Alexandros Kapravelos. Automatic Discovery of Emerging Browser Fingerprinting Techniques. In *Proceedings of the ACM Web Conference 2023*, WWW '23, page 2178–2188, New York, NY, USA, 2023. Association for Computing Machinery.

[57] Karthika Subramani, William Melicher, Oleksii Starov, Phani Vadrevu, and Roberto Perdisci. PhishInPatterns: measuring elicited user interactions at scale on phishing websites. In *Proceedings of the 22nd ACM Internet Measurement Conference*, IMC '22, pages 589–604. Association for Computing Machinery.

[58] Bhaskar Tejaswi, Nayanamana Samarasinghe, Sajjad Pourali, Mohammad Mannan, and Amr Youssef. Leaky Kits: The Increased Risk of Data Exposure from Phishing Kits. In *2022 APWG Symposium on Electronic Crime Research (eCrime)*, pages 1–13, Boston, MA, USA, November 2022. IEEE.

[59] Xiwen Teoh, Yun Lin, Ruofan Liu, Zhiyong Huang, and Jin Song Dong. {PhishDecloaker}: Detecting {CAPTCHA-cloaked} phishing websites via hybrid vision-based interactive models. pages 505–522.

[60] Jonas Tzschoppe and Hans Löhr. Browser-in-the-middle - evaluation of a modern approach to phishing. In *Proceedings of the 16th European Workshop on System Security*, EUROSEC '23, pages 15–20. Association for Computing Machinery.

[61] Rakesh M. Verma and Avisha Das. What's in a url: Fast feature extraction and malicious url detection. *Proceedings of the 3rd ACM on International Workshop on Security And Privacy Analytics*, 2017.

[62] Penghui Zhang, Adam Oest, Haehyun Cho, Zhibo Sun, Rc Johnson, Brad Wardman, Shaown Sarker, Alexandros Kapravelos, Tiffany Bao, Ruoyu Wang, Yan Shoshitaishvili, Adam Doupé, and Gail-Joon Ahn. CrawlPhish: Large-scale Analysis of Client-side Cloaking Techniques in Phishing. In *2021 IEEE Symposium on Security and Privacy (SP)*, pages 1109–1124, San Francisco, CA, USA, May 2021. IEEE.

[63] Penghui Zhang, Adam Oest, Haehyun Cho, Zhibo Sun, RC Johnson, Brad Wardman, Shaown Sarker, Alexandros Kapravelos, Tiffany Bao, Ruoyu Wang, Yan Shoshitaishvili, Adam Doupé, and Gail-Joon Ahn. CrawlPhish: Large-scale analysis of client-side cloaking techniques in phishing. In *2021 IEEE Symposium on Security and Privacy (SP)*, pages 1109–1124. ISSN: 2375-1207.

[64] Penghui Zhang, Zhibo Sun, Sukwha Kyung, Hans Walter Behrens, Zion Leonahenahe Basque, Haehyun Cho, Adam Oest, Ruoyu Wang, Tiffany Bao, Yan Shoshitaishvili, Gail-Joon Ahn, and Adam Doupé. I'm SPARTACUS, No, I'm SPARTACUS: Proactively Protecting Users from Phishing by Intentionally Triggering Cloaking Behavior. In *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security*, CCS '22, pages 3165–3179, New York, NY, USA, November 2022. Association for Computing Machinery.

10 Ethical consideration

This research relied on publicly and commercially available URLs detected using proprietary methods to be phishing websites. No additional threat intelligence is attached to the URLs we are planning to release, and to ensure that any confidential information is not accidentally leaked through the URLs (like GET parameters that are indications of a test submission before the URL was submitted to the feeds), we blind the GET parameters of the URLs we visited upon release.

We did not collect or store any identifiable information about the individuals behind phishing kits or the pages, and we did not conduct any live testing of their networks, as we only visited at most twice upon ingestion.

Appendix A

(a) Page variant in English

(b) Page variant in Japanese

Figure 9: Cropped screenshots from Cluster-e325887b, IP addresses and location redacted to ensure anonymity of the authors.

Table 5: Summary of all API calls that match to a category identified by Crawlphish

Category	Clusters	Pages
User-Interaction		
Pop-up (Total)	104	1,323
Accelerometer	7	63
Clipboard.readText	1	2
Geolocation.getCurrentPosition	55	479
Gyroscope	4	10
MediaDevices.getUserMedia	5	55
Notification.requestPermission	16	148
Window.alert	22	581
Mouse	35	62
DomEvents	10,402	362,700
Fingerprint		
Total	9,131	289,014
Navigator.userAgent	8,641	266,126
HTMLDocument.cookie	4,664	124,110
HTMLDocument.referrer	2,175	39,344
Bot Detection		
Timing	2,528	78,202

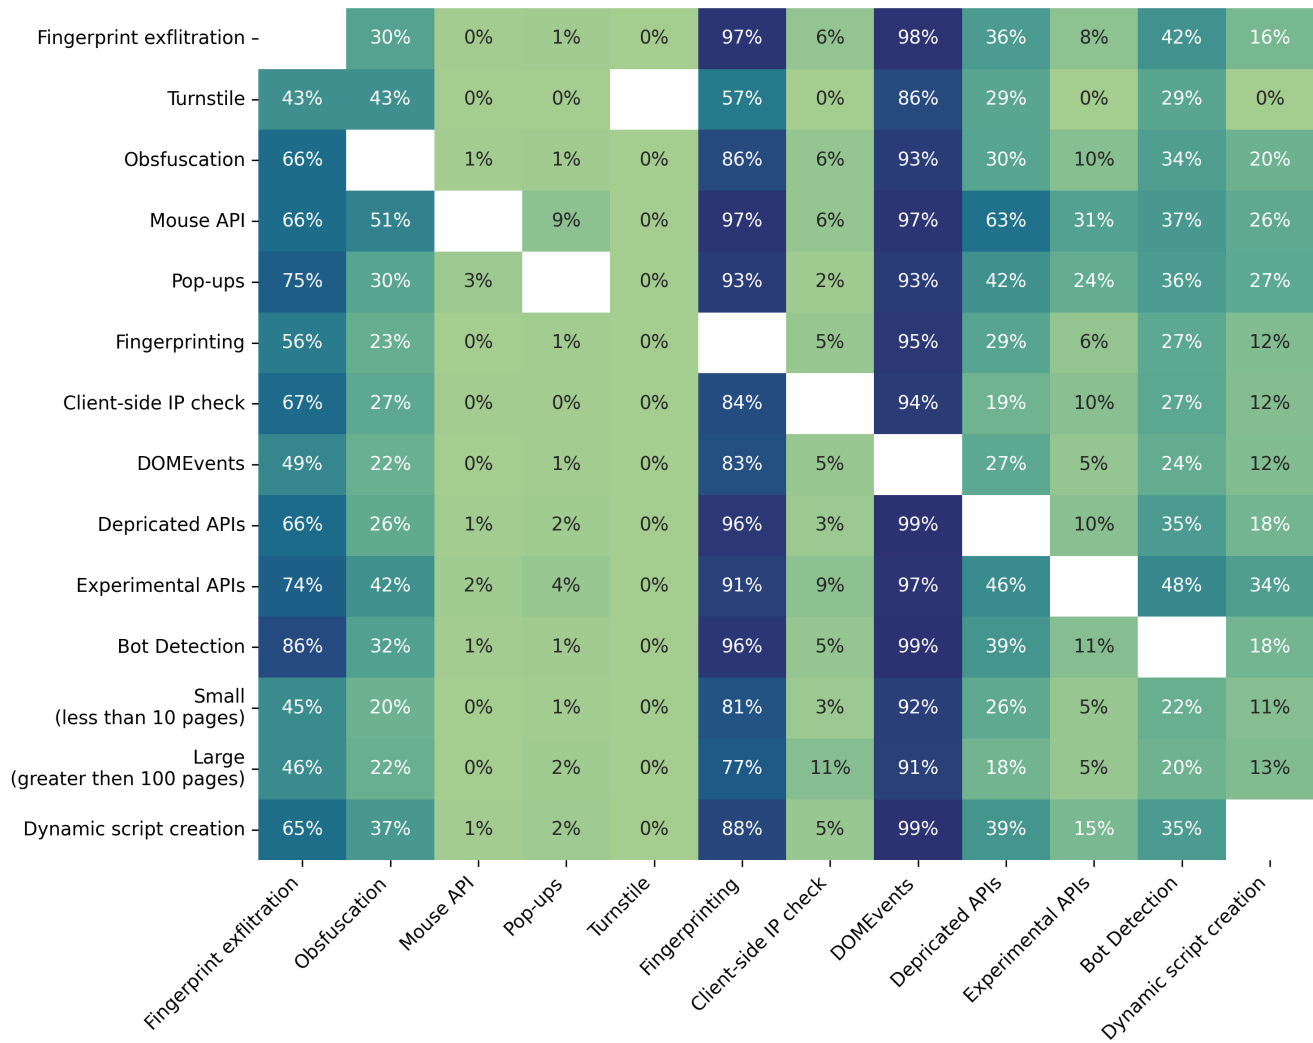|	Fingerprint exflitration	Obsfuscation	Mouse API	Pop-ups	Turnstile	Fingerprinting	Client-side IP check	DOMEvents	Depricated APIs	Experimental APIs	Bot Detection	Dynamic script creation
Fingerprint exflitration		30%	0%	1%	0%	97%	6%	98%	36%	8%	42%	16%
Turnstile	43%	43%	0%	0%		57%	0%	86%	29%	0%	29%	0%
Obsfuscation	66%		1%	1%	0%	86%	6%	93%	30%	10%	34%	20%
Mouse API	66%	51%		9%	0%	97%	6%	97%	63%	31%	37%	26%
Pop-ups	75%	30%	3%		0%	93%	2%	93%	42%	24%	36%	27%
Fingerprinting	56%	23%	0%	1%	0%		5%	95%	29%	6%	27%	12%
Client-side IP check	67%	27%	0%	0%	0%	84%		94%	19%	10%	27%	12%
DOMEvents	49%	22%	0%	1%	0%	83%	5%		27%	5%	24%	12%
Depricated APIs	66%	26%	1%	2%	0%	96%	3%	99%		10%	35%	18%
Experimental APIs	74%	42%	2%	4%	0%	91%	9%	97%	46%		48%	34%
Bot Detection	86%	32%	1%	1%	0%	96%	5%	99%	39%	11%		18%
Small (less than 10 pages)	45%	20%	0%	1%	0%	81%	3%	92%	26%	5%	22%	11%
Large (greater then 100 pages)	46%	22%	0%	2%	0%	77%	11%	91%	18%	5%	20%	13%
Dynamic script creation	65%	37%	1%	2%	0%	88%	5%	99%	39%	15%	35%	

Figure 10: Confusion matrix between all of the techniques enumerated and cluster lifetime characteristics, normalized by row.